

CSIE5428 Computer Vision Practice with Deep Learning


Homework 1 Report

Name: 高榮浩

ID: R12922127

1. The Architecture of Object Detector

I've chosen the DINO model. The diagram below illustrates the DINO framework.



2. Implement Details

I predominantly opt for the original DINO model. The pre-trained checkpoint I leverage from the model zoo is specifically DINO-4scale, with the Swin-L backbone. Consequently, this implementation involves the utilization of two pre-trained weights, namely those for DINO-4scale and Swin-L, as indicated in the following table.

| Name | File | Dataset | Source |
|-----------------------------------|--------------------------------------------|--------------|----------------------|
| DINO-4scale (36 epoch setting) | checkpoint0029_4scale_swin.pth | COCO 2017 | Link |
| Swin-L | swin_large_patch4_ window12_384_22k.pth | ImageNet-22K | Link |

Please note that I have deliberately avoided using the DINO-5scale (36 epoch setting) pre-trained checkpoint. While it does offer a higher box AP, my GPU (GeForce RTX™ 2080 Ti 11G) lacks the capacity to accommodate it. This is due to a `RuntimeError` that occurs, specifically "CUDA out of memory," even with the batch size initially set to 1.

In that case, I attempted to use DINO-4scale (36 epoch setting) with an initial batch size of 2, but my GPU still couldn't handle it. As a result, I had to reduce the batch size to 1, and that ultimately led to successful execution.

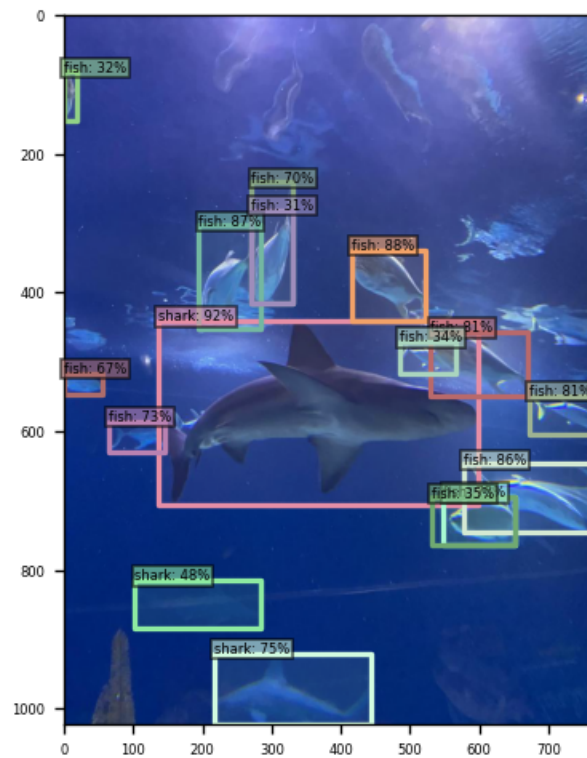
Subsequently, I embarked on a series of experiments involving various combinations of epoch numbers and learning rates, the results of which are outlined in the following section.

3. Performance for Validation Set

| # | Learning Rate | Epoch | Best Epoch | AP | AP ₅₀ | AP ₇₅ |
|----|---------------|-------|------------|--------------|------------------|------------------|
| 1 | 0.0001 | 12 | 11 | 0.582 | 0.863 | 0.608 |
| 2 | 0.0001 | 12 | 11 | 0.574 | 0.850 | 0.610 |
| 3 | 0.0001 | 12 | 6 | 0.577 | 0.858 | 0.611 |
| 4 | 0.0001 | 12 | 11 | 0.578 | 0.852 | 0.609 |
| 5 | 0.0001 | 12 | 11 | 0.574 | 0.863 | 0.587 |
| 6 | 0.0001 | 18 | 12 | 0.578 | 0.848 | 0.606 |
| 7 | 0.0001 | 18 | 17 | 0.580 | 0.866 | 0.605 |
| 8 | 0.0001 | 24 | 16 | 0.587 | 0.863 | 0.609 |
| 9 | 0.0001 | 24 | 18 | 0.581 | 0.859 | 0.604 |
| 10 | 0.0001 | 36 | 19 | 0.584 | 0.864 | 0.618 |
| 11 | 0.00005 | 48 | 36 | 0.588 | 0.866 | 0.621 |
| 12 | 0.000025 | 36 | 23 | 0.587 | 0.867 | 0.596 |
| 13 | 0.000025 | 96 | 12 | 0.591 | 0.867 | 0.616 |
| 14 | 0.0000125 | 36 | 20 | 0.589 | 0.856 | 0.604 |
| 15 | 0.00000625 | 48 | 31 | 0.581 | 0.858 | 0.601 |

This table displays the performance results for the validation set. The highest AP score achieved is 0.591, which was obtained with a learning rate of 0.000025 during the 13th epoch. It's important to note that the index for the 'Best Epoch' in this table begins at 0.

4. Visualization



The image above showcases the detection results for `IMG_2570_jpeg-jpg.rf.ed40900b657a5b23d92cb2d296ad2dbc.jpg` in the testing set.